

Chapter 5 Logging and Recovery

Requirements & Basic Notions,
Logging,
Insertion Strategy (Non-/Atomic),
Propagation Strategy (No-/Force),
Replacement Strategy (No-/Steal),
Savepoints,
Restart-Procedure



Requirements / Basic Notions (1)

- **DBMS Task**
 - Automatic Handling of expectable failures
- **Expectable Failures**
 - DB-Operation rejected
 - Commit not accepted
 - Power breakdown
 - Devices do not work (e.g. magnetic disk)
 - ...
- **Special Characteristics of DBMS Failure Handling**
 - Restriction to and reparation of runtime failures (failure tolerant systems)
 - „Reparation“ of static DB structures



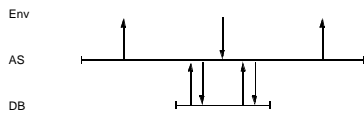
Requirements / Basic Notions (2)

- **General Problems**
 - Failure detection
 - Failure localization
 - Estimation of damage
 - Recovery (itself)
- **Failure Model of centralized DBMS**
 - Transaction failure
 - System failure
 - Device failure
 - Disaster
- **Precondition**
 - Collecting redundant information during normal operation (Logging)



Requirements / Basic Notions (3)

- **Transaction paradigm requires**
 - "All or Nothing" (Atomicity)
 - Durability
- **Goal of Recovery**
 - Most recent transaction-consistent DB state
- **System Environment?**
 - Operating system, application system, other components



Requirements / Basic Notions (4)

- **Basic Forms of Recovery**
 - Forward-Recovery
 - Find a state, at which system can continue to operate
 - However, non-stop paradigm not generally applicable
 - Backward-Recovery
 - Back to most recent consistent state and further processing from there
 - Requires that at all abstractional layers it is clearly defined, to which state it must be restored in case of failure



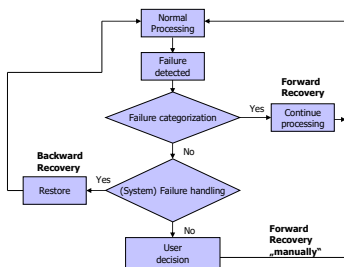
Requirements / Basic Notions (5)

- **Overhead**
 - "A recoverable action is 30% harder and requires 20% more code than a non-recoverable action" (J. Gray)
 - Statement and transaction atomicity required
 - 2 principles
 - **Do things twice:**
prepare first; if OK then concrete modification
 - **Do things once:**
immediate modification; in case of failure internal restore
 - Usually second principle is used (more optimistic and efficient)



Requirements / Basic Notions (6)

Basic processing



Requirements / Basic Notions (7)

Failure classes

Effect	Failure type	Failure class
A single transaction	Violation of system restrictions Security violation Excessive resource requation Application failures e.g. wrong operations or values	Transaction failure
Several transactions	Planned system shut down Problems of resource allocation System overload Deadlock	System failure
All transactions (overall system)	System break down with loss of main memory contents Damage of secondary storage Damage of computer center	System failure Device failure Disaster

Requirements / Basic Notions (8)

Preconditions for Recovery

- quasi-stable storage
- accurate DBMS code
- accurate Log data
- independence of failures

Recovery Classes

1. Transaction Recovery (R1)
- UNDO of single not-committed transaction during database operation (transaction failure, deadlock)
 - Forms
 - Complete UNDO to initial state
 - Partial UNDO to savepoint within transaction

Requirements / Basic Notions (9)

- **Recovery Classes (contd.)**

- 2. **Crash Recovery** (R2) after System Crash

- Restore of most recent transaction consistent DB state
 - Necessary actions
 - (partial) REDO of successful transactions (REDO of lost modifications)
 - UNDO of all interrupted transactions (removal of all their modifications from permanent DB)

- 3. **Media Recovery** (R3) after Device Failure

- Mirroring (at disk level)
 - Complete REDO of all modifications of successful completed transactions on archive copy of DB



Requirements / Basic Notions (10)

- **Recovery Classes (contd.)**

- 4. **Disaster Recovery** (R4)

- DB copy in remote system
 - Delayed continuation of DB processing on repaired/new system on basis of archive copy (possibly data loss)



Requirements / Basic Notions (11)

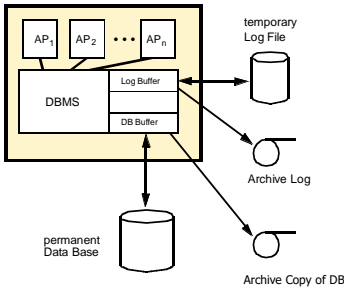
- **Not (formally) classified**

- R5 Recovery
 - Log-Data damaged
 - R6 Recovery
 - Beyond DBMS
 - Compensation transactions
 - Manual treatment



Requirements / Basic Notions (12)

DB-Recovery – System Components



Requirements / Basic Notions (13)

DB-Recovery – System Components (contd.)

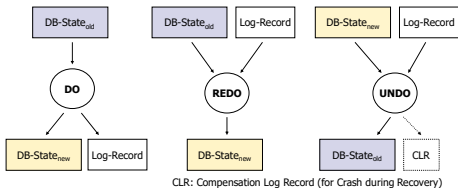
- Buffering Log Data in Main Memory (Log Buffer)
 - Propagation (at the latest) at Commit
- Usage of Log Data
 - Temporary Log File for Handling Transaction Failures and System Failures
 - DB + temp. Log ⇒ DB
 - Handling Device Failures
 - Archive Copy + Archive Log ⇒ DB

Logging (1)

Task

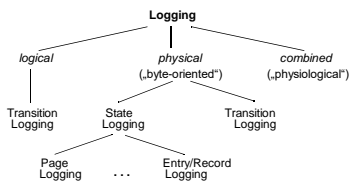
- Collecting redundant data w.r.t. modifications during normal DB processing
- Usage in case of failure (Undo-, Redo-Recovery)

Do-Redo-Undo-Principle



Logging (2)

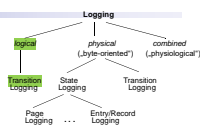
- Logging can be performed at (DBMS-) System Layer
- Logging Techniques:



Logging (3)

Logical Logging

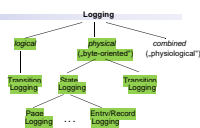
- Logging update DML operations with parameters
- General Problem
 - Set oriented update operations
- Precondition
 - After system crash persistent data base must at least be action consistent, in order to being able to perform „reverse operations“
- Deferred (indirect) insertion strategy needed



Logging (4)

Physical Logging

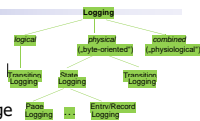
- Log Granulate: Page vs. Entry/Record
- State Logging
 - Before Images and After-Images are stored in Log File
- Transition Logging
 - Difference between Before- and After-Image is stored
- Applicable with direct as well as indirect insertion strategies
- Problems of logical and physical Logging Techniques**
 - Logical Logging: not compatible to Update-in-Place
 - Physical, „byte-oriented“ Logging: complex and inflexible w.r.t. deletion and insertion operations



Logging (5)

Physiological Logging

- Combined physical/logical Logging
 - physical-to-a-page, logical-within-a-page
 - Each Log-Record relates to a single database page
 - Logging of elementary, internal operations within a page
 - Compatible to Update-in-Place



Logging (6)

Examples

- Modifications of a page A
 1. Inserting object a into page A ($A_1 \rightarrow A_2$)
 2. Modifying object b_{old} to b_{new} in page A ($A_2 \rightarrow A_3$)

	logical	physical
States		Logging Before- and After-Images 1. A_1 and A_2 2. A_2 and A_3
Transitions	Logging Operations with Parameters 1. Insert (a) 2. Update (b_{old}, b_{new})	Logging 'Differences' 1. $A_1 \oplus A_2$ 2. $A_2 \oplus A_3$

Logging (7)

Examples (contd.)

- Logging „Diffs“: Reconstruction of Pages
 - A1 as starting point or A3 as endpoint available
 - REDO-Recovery
 - $A_1 \oplus (A_1 \oplus A_2) = A_2$
 - $A_2 \oplus (A_2 \oplus A_3) = A_3$
 - UNDO-Recovery
 - $A_3 \oplus (A_2 \oplus A_3) = A_2$
 - $A_2 \oplus (A_1 \oplus A_2) = A_1$

A	B	XOR
0	0	0
0	1	1
1	0	1
1	1	0

Logging (8)

Assessment of Logging Techniques

	Overhead during normal processing	Restart-Overhead after failure (Crash)
Page-Logging	--	+
Page-Transition-Logging (Differences)	-	+
Entry/Record-Logging / physiological Logging	+	+
Logical Logging	++	--



Logging (9)

Entry-Logging vs. Page-Logging

- Advantages of Entry-Logging
 - Low storage overhead
 - Less Log-I/Os
 - Allows 'better' buffering of Log data (Group-Commit)
 - Supports more fine-grained concurrency control granulates (Page-Logging → CC on page level)
- Drawback of Entry-Logging
 - Recovery more complex than in case of Page-Logging
 - e.g. before application of log records pages must be loaded into main memory



Logging (10)

Structure of (temporary) Log-File

- Several different types of Log Records needed
 - BOT-, Commit-, Abort-Records
 - Update-Record (UNDO-Information, e. g. 'Before-Images', and REDO-Information, e. g. 'After-Images')
 - Checkpoint-Records
- Update Record
 - Structure of Record: [LSN, TAID, PageID, Redo, Undo, PrevLSN]
 - LSN: Log Sequence Number
 - Unique ID of log record
 - LSNs are created in monotonically ascending order
 - Thus, chronological order of logging entries can be reconstructed



Logging (11)

Structure of (temporary) Log-File (contd.)

- TAID: ID of transaction which issued update
- PageID
 - ID of modified page
 - If modification relates to more than one page then several log records need to be created
- Redo
 - Redo-Information specifies how modification can be reproduced
- Undo
 - Undo-Information specifies how modification can be withdrawn
- PrevLSN
 - Pointer to previous log record of the same transaction
 - Needed for efficiency reasons during transaction UNDO



Logging (12)

Structure of (temporary) Log-File (contd.)

Example

Schritt	T ₁	T ₂	Log
			[LSN, TAID, PageID, Redo, Undo, PrevLSN]
1.	BOT		[#1, T ₁ , BOT, 0]
2.	r(A, a ₁)		
3.		BOT	[#2, T ₂ , BOT, 0]
4.		r(C, c ₂)	
5.	a ₁ := a ₁ - 50		
6.	w(A, a ₂)		[#3, T ₁ , P _A , A=50, A+=50, #1]
7.		c ₂ := c ₂ + 100	
8.		w(C, c ₂)	[#4, T ₂ , P _C , C=100, C+=100, #2]
9.	r(B, b ₁)		
10.	b ₁ := b ₁ + 50		
11.	w(B, b ₁)		[#5, T ₁ , P _B , B+=50, B+=50, #3]
12.	Commit		[#6, T ₁ , Commit, #5]
13.	r(A, a ₂)		
14.	a ₂ := a ₂ - 100		
15.	w(A, a ₂)		[#7, T ₂ , P _A , A=100, A+=100, #4]
16.		Commit	[#8, T ₂ , Commit, #7]



Logging (13)

Structure of (temporary) Log-File (contd.)

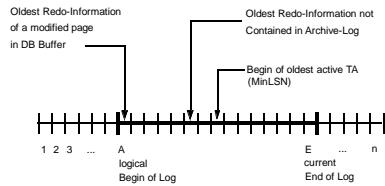
- Sequential File
 - Writing new logging data at end of file
- Log-Data relevant for Crash Recovery only for restricted period of time
 - Undo-Information no longer needed as soon as transaction is completed successfully
 - After insertion of page into permanent DB Redo-Information is no longer needed
 - Redo-Information for Media-Recovery is collected in Archive-Log!



Logging (14)

Structure of (temporary) Log-File (contd.)

- Ring Buffer



Related System Components (1)

Overview

- Insertion Strategy
 - Direct insertion of modifications into permanent DB (*non-atomic*)
 - Deferred insertion (*atomic*)
- Replacement Strategy
 - Pushing 'dirty' pages to secondary storage (*steal*)
 - Only pages of successfully completed transactions are pushed (*nosteal*)
- Propagation Strategy
 - Propagation at Commit mandatory (*force*)
 - Propagation possibly after Commit (*noforce*)
- Locking Granulate
- Commit-Procedure

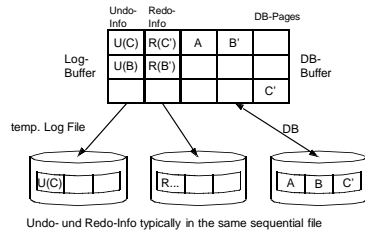
Related System Components (2)

Insertion Strategy

- non-atomic / direct / update-in-place*
 - Modified page is always stored back to the same block on disk
 - 'storing back' (here) means insertion into permanent DB
 - 'atomic' insertion of several pages not possible (non-atomic)
- Requirements
 - WAL-Principle: *Write Ahead Log* for Undo-Info; U(B) before B' (cf. next slides)
 - Logging Redo-Info at the latest at Commit; R(C) + R(B') before Commit (cf. next slides)

Related System Components (3)

- **Insertion Strategy (contd.)**
 - *non-atomic / direct / update-in-place (contd.)*

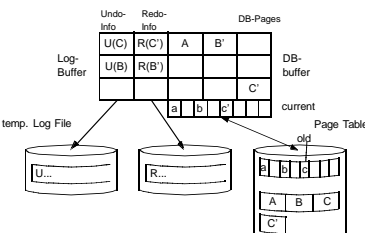


Related System Components (4)

- **Insertion Strategy (contd.)**
 - *atomic / deferred*
 - e. g. in System R, SQL/DS
 - Modified page is written to separate block on disk, actual 'insertion into DB' is performed later on
 - Page table contains page address
 - Deferred, atomic insertion of multiple modifications possible by switching page tables
 - Action consistent or even transaction consistent DB on disk
 - Thus, logical logging applicable
 - Requirements
 - WAL-Principle:
U(C) + U(B) before checkpoint
 - R(C') + R(B') at the latest at Commit

Related System Components (5)

- **Insertion Strategy (contd.)**
 - *atomic / deferred (contd.)*



Related System Components (6)

Replacement Strategy

- Problem: Replacement of ,dirty' pages
- steal
 - Modified pages can be replaced (in buffer) and stored into the permanent DB at any time, esp. before commit of corresponding TA
 - Higher flexibility for page replacement
 - Undo-Recovery needed (TA Abort, System Crash)
 - steal requires observation of WAL principle, i.e., before writing a dirty page corresponding UNDO Information (e.g. Before Image) must be written into Log File
- nosteal
 - Dirty pages must not be replaced
 - No UNDO Recovery needed
 - Problems in case of long Update TA



Related System Components (7)

Propagation Strategy

- force
 - All modified pages are propagated to permanent DB at the latest at Commit ('writing through')
 - No Redo Recovery needed after System Crash
 - High Overhead
 - Large DB Buffers possibly not really exploited
 - Longer answering times for update TA
- noforce
 - No 'write through' at Commit
 - At (the latest at) Commit only Redo Information must be written to Log File
 - Redo Recovery after System Crash
- Commit-Rule
 - before TA Commit sufficient Redo Information (e. g. After Images) must be written for all modifications



Related System Components (8)

Consequences

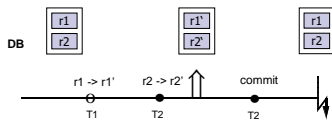
	steal	nosteal
force	UNDO NO REDO	NO UNDO NO REDO
noforce	UNDO REDO	NO UNDO REDO



Related System Components (9)

▪ Lock Management

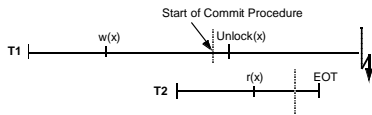
- Log Granulate must be smaller as or equal to Lock Granulate
 - Example
 - Record Locks
 - Before and After Images at page level
 - Undo (Redo) of a modification can overwrite parallel modification of the same page (*lost update*)



Related System Components (10)

▪ Commit Procedure

- Requirements
 - Modifications need to be „assured“ at Commit
 - Modifications get visible to other TA not before it can be assured that corresponding update TA will get to its Commit (Problem of recursive Aborts)



Related System Components (11)

▪ Commit Procedure (contd.)

- 2-phase processing
 - Phase 1: ensuring repeatability of TA
 - Storing modifications
 - Writing Commit Record to Log
 - Phase 2: making modifications visible to others (Releasing Locks)
 - At the end of phase 1 user/application can be informed that TA has been successful
- Example: Commit Procedure for *force, steal*:
 1. Writing Before Images to Log
 2. Force of modified DB Pages
 3. Writing After-Images (for Archive Log) and Commit RecordIn case of NoForce just 3. needed for first Commit Phase



Related System Components (12)

Commit Procedure (contd.)

- Group Commit
 - Log File is potential bottleneck
 - At least 1 Log I/O for each update TA
 - max. about 250 sequential writes per second (1 disk)
 - Group Commit means writing Log Data of several TA
 - Buffering Log Data in Log Buffer (1 or more pages)
 - Precondition: Record Logging
 - Log Buffer is written to Log File if full or time limit exceeded (Timer)
 - Insignificant delay of Commit



Related System Components (13)

Commit Procedure (contd.)

- Group Commit (contd.)
 - Group Commit allows reduction to 0.1 - 0.2 Log-I/Os per TA
 - Less CPU overhead (for I/Os) reduces waiting times for CPU
 - Dynamic adjustment of timer value by DBMS desirable
 - Thus, Group Commit allows increase of throughput, esp. w.r.t. log bottleneck and high CPU utilization



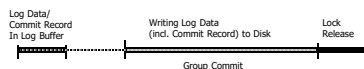
Related System Components (14)

Commit Procedure (contd.)

- Comparison
 - Standard 2PC



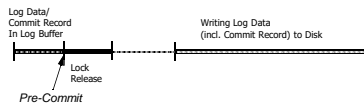
- Group Commit



Related System Components (15)

Commit Procedure (contd.)

- Further Optimization: *Pre-Commit*
 - Lock release after Commit Record has been written to Log Buffer (not Log File)
 - TA can only be aborted by system crash
 - In this case all depending TAs (having seen 'unstored' modifications because of early lock release) fail too

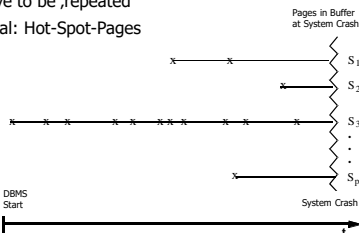


- In all 3 variants user is informed about TA termination not before Commit Record has been written to external storage (temp. log file)

Checkpoints (1)

Motivation: Checkpoints

- Goal: restricting REDO overhead after System Crash (noforce)
- Without checkpoint potentially all modifications since system start would have to be 'repeated'
- Especially critical: Hot-Spot-Pages



Checkpoints (2)

Management Data

- Log File
 - BEGIN_CHKPT Record
 - (actual) Checkpoint Information, e. g. list of active TAs
 - END_CHKPT Record
- Log Address of last Checkpoint Record is kept in special System File

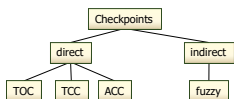
Checkpoints and 'Insertion'

- atomic*
 - State of permanent DB is state of last (successful) checkpoint
- non-atomic*
 - State of permanent DB contains all pages inserted before crash

Checkpoints (3)

Kinds of Checkpoints

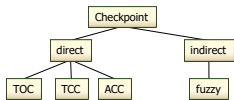
- Direct Checkpoints
 - All modified pages are stored into the permanent DB
 - REDO Recovery starts at last Checkpoint
 - Drawback: long system down time, because there must be no modifications 'during Checkpoint'
 - *Transaction consistent or action consistent Checkpoints*



Checkpoints (4)

Kinds of Checkpoints (contd.)

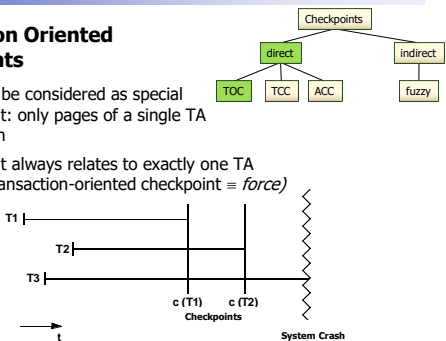
- Indirect/Fuzzy Checkpoints
 - Modified pages do not have to be propagated
 - Only state information (IDs of pages in Buffer, active TAs, open files etc.) are written to Log File
 - Minor (Checkpoint-) Overhead
 - Generally, REDO Information before checkpoint must be taken into account
 - Special treatment of Hot-Spot-Pages



Checkpoints (5)

Transaction Oriented Checkpoints

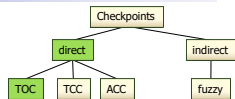
- Force can be considered as special Checkpoint: only pages of a single TA are written
- Checkpoint always relates to exactly one TA (TOC = transaction-oriented checkpoint = *force*)



Checkpoints (6)

Transaction Oriented Checkpoints (contd.)

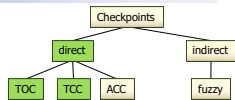
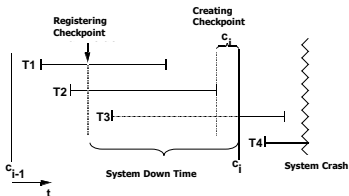
- Properties
 - Commit processing enforces propagation of all modified pages of the TA
 - Insertion of all modifications into permanent DB
 - Comment in Log File
 - only *atomic* supports atomic insertion of multiple pages
 - Thus, at least in case of direct Insertion of pages UNDO Recovery has to be provided (*steal*)



Checkpoints (7)

Transaction Consistent Checkpoints

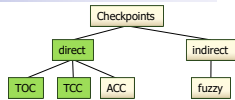
- Checkpoint always relates to all TA (TCC = transaction consistent checkpoint)



Checkpoints (8)

Transaction Consistent Checkpoints (contd.)

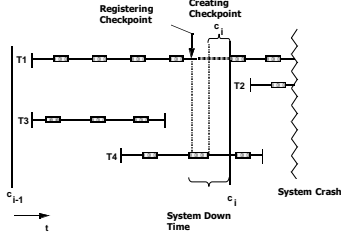
- Properties
 - Propagation to be deferred until end of all active update TA
 - New update TA must wait, until checkpoint creation completed
 - Crash Recovery starts at last checkpoint



Checkpoints (9)

Action Consistent Checkpoints

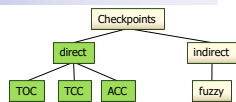
- Checkpoint always relates to all TA (ACC = action consistent checkpoint)



Checkpoints (10)

Action Consistent Checkpoints (contd.)

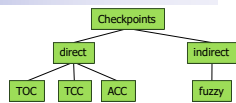
- Properties
 - No update statements during checkpoint
 - In comparison to TCC shorter down times, but lower checkpoint 'quality'
 - Crash Recovery not limited by last checkpoint



Checkpoints (11)

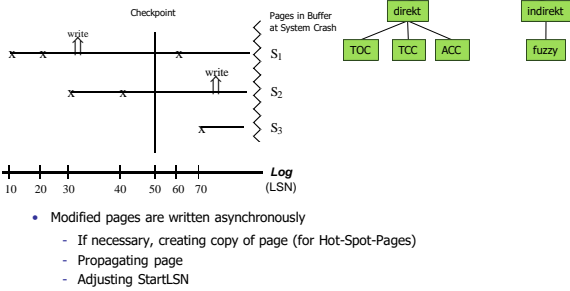
Fuzzy Checkpoints

- DB stays 'fuzzy'
 - Only relevant for Update-in-Place (*non-atomic*)
- Problem
 - Determination of Log position, at which Redo Recovery has to start
 - Buffer manager stores StartLSN for each modified page, i.e., LSN of first modification since reading from disk
 - Redo-Recovery after Crash starts at MinDirtyPageLSN (= MIN(StartLSN))
- Checkpoint Information
 - MinDirtyPageLSN, List of active TA and their StartLSNs, ...



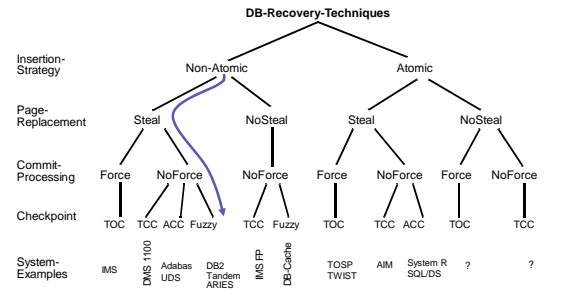
Checkpoints (12)

Fuzzy Checkpoints (contd.)



Checkpoints (13)

Combinations



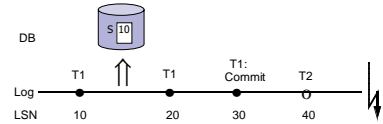
Recovery (1)

LSNs

- Challenge
 - Decision at restart, whether or not Recovery action is necessary for considered page (old or new state on external storage?)
 - For that purpose each page header contains LSN of most recent Log Entry L related to this page (PageLSN (B) := LSN (L))
- Decision Procedure
 - Restart contains Redo- and Undo-Phase
 - Redo necessary, only if $\text{Page-LSN} < \text{LSN of Redo-Log-Entry}$
 - Undo necessary, only if $\text{Page-LSN} \geq \text{LSN of Undo-Log-Entry}$

Recovery (2)

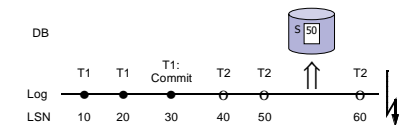
- LSNs (contd.)
- Simplified application: Page Locks



- Redo T1: $S(10) = T1(10): -$
 $S(10) < T1(20): \text{Redo, } S(20)$
- Undo T2: $S(20) < T2(40): -$
- Page-LSN is updated at Redo (increases monotonically)

Recovery (3)

- LSNs (contd.)
- Simplified application: Page Locks (contd.)



- Redo T1: $S(50) > T1(10): -$
 $S(50) > T1(20): -$
- Undo T2: $S(50) < T2(60): -$
 $S(50) \geq T2(50): \text{Undo}$
 $S(50) \geq T2(40): \text{Undo}$

Recovery (4)

- LSNs (contd.)
- Undo in LIFO Order
- Special treatment of „UNDOS“ necessary so that repeated application leads to the same result (idempotent)
- State Logging and LIFO Order ensure that processing is idempotent

Recovery (5)

Crash-Recovery

- Goal
 - Creating the most recent transaction consistent DB state from permanent DB and temporary Log File
- In case of Update-in-Place (*non-atomic*)
 - State of permanent DB after Crash unpredictable („chaotic“)
 - Thus, only physical (or physiological) Logging applicable
 - A Block of the permanent DB either is
 - Up-to-date
 - or outdated (*noforce*) → Redo
 - or „dirty“ (*steal*) → Undo



Recovery (6)

Crash-Recovery (Forts.)

- In case of *atomic*
 - State of permanent DB corresponds to the most recent successful propagation (checkpoint)
 - At least action consistent → DML statements can be executed (logical Logging)
 - *force*: no Redo
 - *noforce*:
 - Transaction consistent propagation → Redo, no Undo
 - Action consistent propagation → Undo + Redo



Recovery (7)

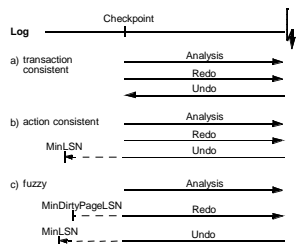
General Restart Procedure

- 3 Phases
 1. Analysis Phase
 - From last checkpoint to end of Log
 - Determination of winner and loser TAs as well as of modified pages
 2. Redo Phase
 - Reading Log forward: starting point depends on checkpoint type
 - selective Redo (*redo winners*) in case of page locks or complete Redo (*repeating history*)
 3. Undo Phase
 - UNDO of all „losers“
 - Reading Log backward until BOT record of oldest loser TA



Recovery (8)

General Restart Procedure (contd.)



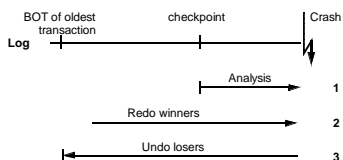
Recovery (9)

Restart Procedure (Update-in-Place)

- Properties: non-atomic, steal, noforce, fuzzy checkpoints
- Process
 1. Analysis Phase
 - From last checkpoint to end of log
 2. Redo Phase
 - Starting point depends on checkpoint type: here MinDirtyPageLSN
 - selective Redo: modifications of winner TAs only
 3. Undo Phase
 - Loser TA up to MinLSN

Recovery (10)

Restart Procedure (Update-in-Place) (contd.)



- Overhead
 - For steps 2 and 3 corresponding pages must be loaded from external storage
 - Page LSNs indicate, whether or not Log information must be applied
 - At the end all modified pages must be propagated again, or a checkpoint is created, respectively

Recovery (11)

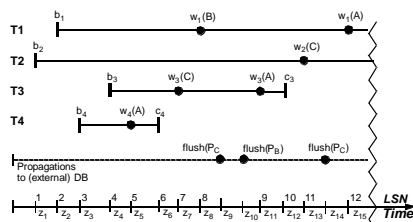
▪ Redo

- In case of physical and physiological Logging
 - Redo action for Log record L is determined by PageLSN of corresponding page B

```
if (B not in Buffer) then
  load B into main memory;
fi;
if (LSN (L) > PageLSN (B)) then
  Redo (modification from L);
PageLSN (B) := LSN (L);
fi;
```
- Repeated application of Log record (e.g. after multiple failures) keeps correctness (REDO is idempotent)
- Recovery in case of Crash during Restart?

Recovery (12)

▪ Restart Example:



Recovery (13)

▪ Restart Example (contd.)

- Assumptions
 - At the beginning: all Page-LSNs 0
- Analysis Phase:
 - Winner-TA: T3, T4
 - Loser-TA: T1, T2
 - relevant pages: P_A, P_B, P_C
- Comment
 - In the example: page locks
 - Thus, selective Redo sufficient (Redo only for winners)

Recovery (14)

Restart Example (contd.)

- Redo Phase:
 - Checking Log records for T₃ and T₄ (forward)

TA	Page	Page-LSN	Log-Record-LSN	Action
T ₄	P _A	0 → 5	5	REDO
T ₃	P _C	11	7	no REDO
T ₃	P _A	5 → 9	9	REDO

- Redo only, if Page-LSN < Log-Record-LSN
- Page-LSNs increase monotonically



Recovery (15)

Restart Example (contd.)

- Undo Phase:
 - Checking Log records for T₁ and T₂ (backward)

TA	Page	Page-LSN	Log-Record-LSN	Action
T ₁	P _A	9	12	no Undo
T ₂	P _C	11	11	Undo
T ₁	P _B	8	8	Undo

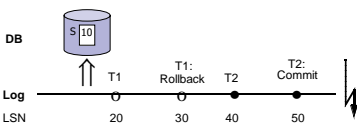
- Undo only, if Page-LSN ≥ Log-Record-LSN
- Because of page logs there is no interference between REDO and UNDO actions; state logging ensures that UNDO is idempotent



Recovery (16)

UNDO Problems w.r.t. LSN Exploitation

- Problem 1: TA UNDO
 - Taking previous Rollbacks into account?



- Redo of T2: S(10) < T2(40) : Redo, S(40)
- Undo of T1: S(40) > T1(20) : Undo, Failure



Recovery (17)

UNDO Problems w.r.t. LSN Exploitation (contd.)

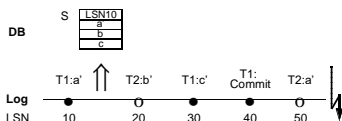
- Problem 1: TA UNDO (contd.)
 - Comment
 - UNDO of modification 20, although modification not represented by page S
 - Assigning LSN = 20 to S violates monotonicity requirement for Page LSNs



Recovery (18)

UNDO Problems w.r.t. LSN Exploitation (contd.)

- Problem 2: Record Locks
 - T1 and T2 modify page S concurrently



- Redo T1: $S(10) \geq T1(10)$: no Redo
 $S(10) < T1(30)$: Redo, S(30)
- Undo T2 (LIFO): $S(30) < T2(50)$: no Undo
 $S(30) > T2(20)$: Undo, Failure!
- More general UNDO processing needed!



Recovery (19)

Failure tolerance of Restart

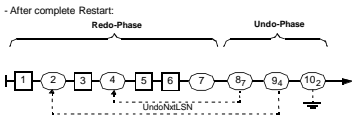
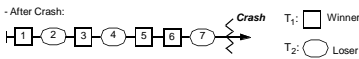
- Requirement: Restart must be idempotent
$$\text{Undo}(\text{Undo}(\dots (\text{Undo}(A)) \dots)) = \text{Undo}(A)$$
$$\text{Redo}(\text{Redo}(\dots (\text{Redo}(A)) \dots)) = \text{Redo}(A)$$
- Solution
 - REDO idempotent since Page LSNs increase monotonically
 - „Compensation Log Records“ for UNDO



Recovery (21)

Failure tolerance of Restart

- CLR = Compensation Log Record
 - CLRs for all Undo actions: Rollback and Undo Phase
 - in Redo Phase: complete Redo of Winners and Losers („repeating history“)
 - Illustration of Log File



Recovery (22)

Failure tolerance of Restart (contd.)

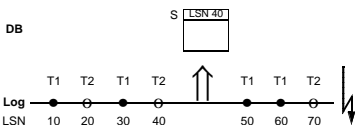
- CLR = Compensation Log Record (contd.)
 - Redo Information of CLR equates to UNDO operation as performed in UNDO phase
 - CLRs are needed for repeated Restart (Crash during Restart); then their REDO Information is applied and corresponding Page LSNs are modified → idempotent
 - CLRs do not need Undo Information; they are skipped in subsequent Undo Phases (UndoNxtLSN)



Recovery (23)

Compensation Log Records (contd.)

- Example
 - all modifications relate to page S
 - State after Crash 1



Recovery (24)

Compensation Log Records (contd.)

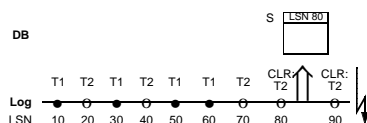
- Example (contd.)
 - State after Crash 1 (contd.)
 - repeating history:
 - $S(40) > T1(10)$: -
 - ...
 - $S(40) \geq T2(40)$: -
 - $S(40) < T1(50)$: Redo, $S(50)$
 - $S(50) < T1(60)$: Redo, $S(60)$
 - $S(60) < T2(70)$: Redo, $S(70)$
 - Undo of T2:
 - CLR(80): Compensation of T2(70), S(80)
 - Propagating S to DB (Flush S)
 - CLR(90): Compensation of T2(40), S(90)
 - Crash



Recovery (25)

Compensation Log Records (contd.)

- Example (contd.)
 - State after Crash 2



Recovery (26)

Compensation Log Records (contd.)

- Example (contd.)
 - State after Crash 2 (contd.)
 - Repeating History:
 - $S(80) > T1(10)$: -
 - ...
 - $S(80) > T2(70)$: -
 - CLR(80): -
 - CLR(90): Compensation of T2(40), S(90)
 - Undo of T2:
 - CLR(100): Compensation of T2(20), S(100)
 - End



Recovery (27)

- **Restart Procedure (Update-in-Place)**
 - Properties: non-atomic, steal, noforce, fuzzy checkpoints
 - 1. Analysis Phase
 - From last checkpoint to end of log
 - 2. Redo Phase
 - Starting point: MinDirtyPageLSN
 - Selective Redo or Repeating History (if necessary)
 - 3. Undo Phase
 - UNDO of losers back to MinLSN



Recovery (28)

- **Restart Procedure (Update-in-Place) (contd.)**
 - Properties: non-atomic, steal, noforce, fuzzy checkpoints (contd.)
 - The diagram shows a horizontal timeline. It starts with 'Log', followed by 'BOT of oldest transaction', then 'checkpoint', and finally 'Crash' indicated by a lightning bolt. Below the timeline, three phases are shown with arrows: 'Analysis' (from checkpoint to end of log), 'Redo winners' (from BOT of oldest transaction to end of log), and 'Undo losers' (from end of log back to BOT of oldest transaction). These are numbered 1, 2, and 3 respectively.
- ARIES
 - Algorithm for Recovery and Isolation Exploiting Semantics
 - Developed by C. Mohan et al. (IBM Almaden Research)
 - Realized in several commercial DBMS

Mohan, C. et al.: A Transaction Recovery Method Supporting Fine-Granularity Locking and Partial Rollbacks Using Write-Ahead Logging, in ACM TODS 17:1, 1992, 94-162



C. Mohan, IBM Fellow

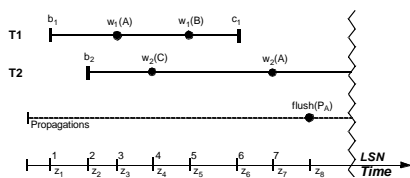


Mohan, C. et al.: A Transaction Recovery Method Supporting Fine-Granularity Locking and Partial Rollbacks Using Write-Ahead Logging, in ACM TODS 17:1, 1992, 94-162



Recovery (29)

Restart Example 2



Recovery (30)

Restart Example 2 (contd.)

- Analysis Phase
 - Winner-TA: T_1
 - Loser-TA: T_2
 - relevant Pages: P_A, P_B, P_C
- Comment
 - Complete Redo
 - For Undo operations: CLR with the following structure:
[LSN, TAID, PageID, Redo, PrevLSN, UndoNextLSN]

Recovery (31)

Restart Example 2 (contd.)

- Redo Phase
 - Checking log records of all TA (T_1, T_2) forwards

TA	Page	Page-LSN	Log-Record-LSN	Action
T_1	P_A	7	3	No REDO
T_2	P_C	$0 \rightarrow 4$	4	REDO
T_1	P_B	$0 \rightarrow 5$	5	REDO
T_2	P_A	7	7	No REDO

- Redo, if Page-LSN < Log-Record-LSN

Recovery (32)

Restart Example 2 (contd.)

- Undo Phase
 - Checking log records of loser TA T_2 backward
 - For each log record Undo is performed and CLR written to log end

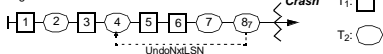
TA	Log-Record-LSN	Action
T_2	7	UNDO and CLR[8, T_2 , P_A , U(A), 7, 4]
T_2	4	UNDO and CLR[9, T_2 , P_C , U(C), 8, 2]
T_2	2	UNDO and CLR[10, T_2 , ..., 9, 0]

Recovery (33)

Restart Example 2 (contd.)

- Assumption
 - Crash during Restart

Log File after Crash:



- Analysis Phase
 - As known
- Redo Phase
 - Checking log records of all TA (T_1 , T_2) incl. CLR forward
 - Redo for each CLR

Recovery (34)

Restart Example 2 (contd.)

- Redo Phase (contd.)

TA	Page	Page-LSN	Log-Record-LSN	Action
T_1	P_A	7	3	No REDO
T_2	P_C	4	4	No REDO
T_1	P_B	5	5	No REDO
T_2	P_A	7	7	No REDO
T_2	P_A	7 → 8	8	REDO: U(A)

Recovery (35)

Restart Example 2 (contd.)

- Undo Phase
 - Checking log records of loser TA T_2 backward
 - For each log record Undo is performed and CLR written to log end

TA	Log-Record-LSN	Action
T_2	8	UndoNxtLSN = 4, go to Log Record 4 (Log Record 7 is skipped, since it is already compensated by 8)
T_2	4	UNDO and CLR[9, T_2 , P_C , U(C), 8, 2]
T_2	2	UNDO and CLR[10, T_2 , r , r , 9, 0]



Conclusion (1)

- Failures
 - Transaction-, System-, Device Failures and Disasters
- Spectrum of Logging- and Recovery-Mechanisms
 - Entry-Logging outmatches Page-Logging
 - Many DBMS use physiological Logging
 - More flexible recovery within a page
 - Less storage overhead
 - Less I/Os
 - Group Commit



Conclusion (2)

- Dependencies to other Components
 - Lock granulate must be greater or equal to log granulate
- Atomic
 - Saves DB state of last checkpoint
 - Ensure action consistency
 - Allow logical logging
- Update-in-Place
 - More effective w.r.t. normal operation
 - Low crash probability
 - Require physical logging



Conclusion (3)

- **Basics w.r.t. Update-in-Place**
 - WAL principle: Write Ahead Log for Undo Info
 - Redo Info to be written at the latest at Commit
- **Basics w.r.t. Atomic**
 - WAL principle:
 - TA-related Undo-Info must be written before checkpoint
 - Redo Info to be written at the latest at Commit
- **NoForce**
 - Outmatches force
 - Require checkpoints in order to limit Redo overhead
 - Fuzzy checkpoints cause lowest overhead w.r.t. normal operation



Conclusion (4)

- **Steal**
 - Requires WAL principle
 - Requires Undo actions after crash
- **Restart**
 - Redo action increase page LSNs
 - CLRs for Undo and Rollback actions
 - Restart idempotent